

Matlab on GPUs

Jerry Ebalunode

Senior Researcher

HPE DSI

uh.edu/datascience

Objective

- Learn how to run MATLAB codes on NVIDIA CUDA-enabled GPUs

Overview

- Review
 - Basics
 - Parallel MATLAB
- MATLAB on GPUs
 - Built-in functions
 - Elementwise operations
 - gpuArrays
 - GPU Code porting
- Conclusion

Review

- MATLAB Introduction

- Variables, arrays, matrices, etc

- Arrays are the fundamental units of data

- Operators `>>sum(M)`

```
>>a= [1 2 3; 4 5 6; 7 8 9; 10 11 12];  
      Column major
```

- Flow control


- if, while, for

```
for index = values  
    statements  
end
```

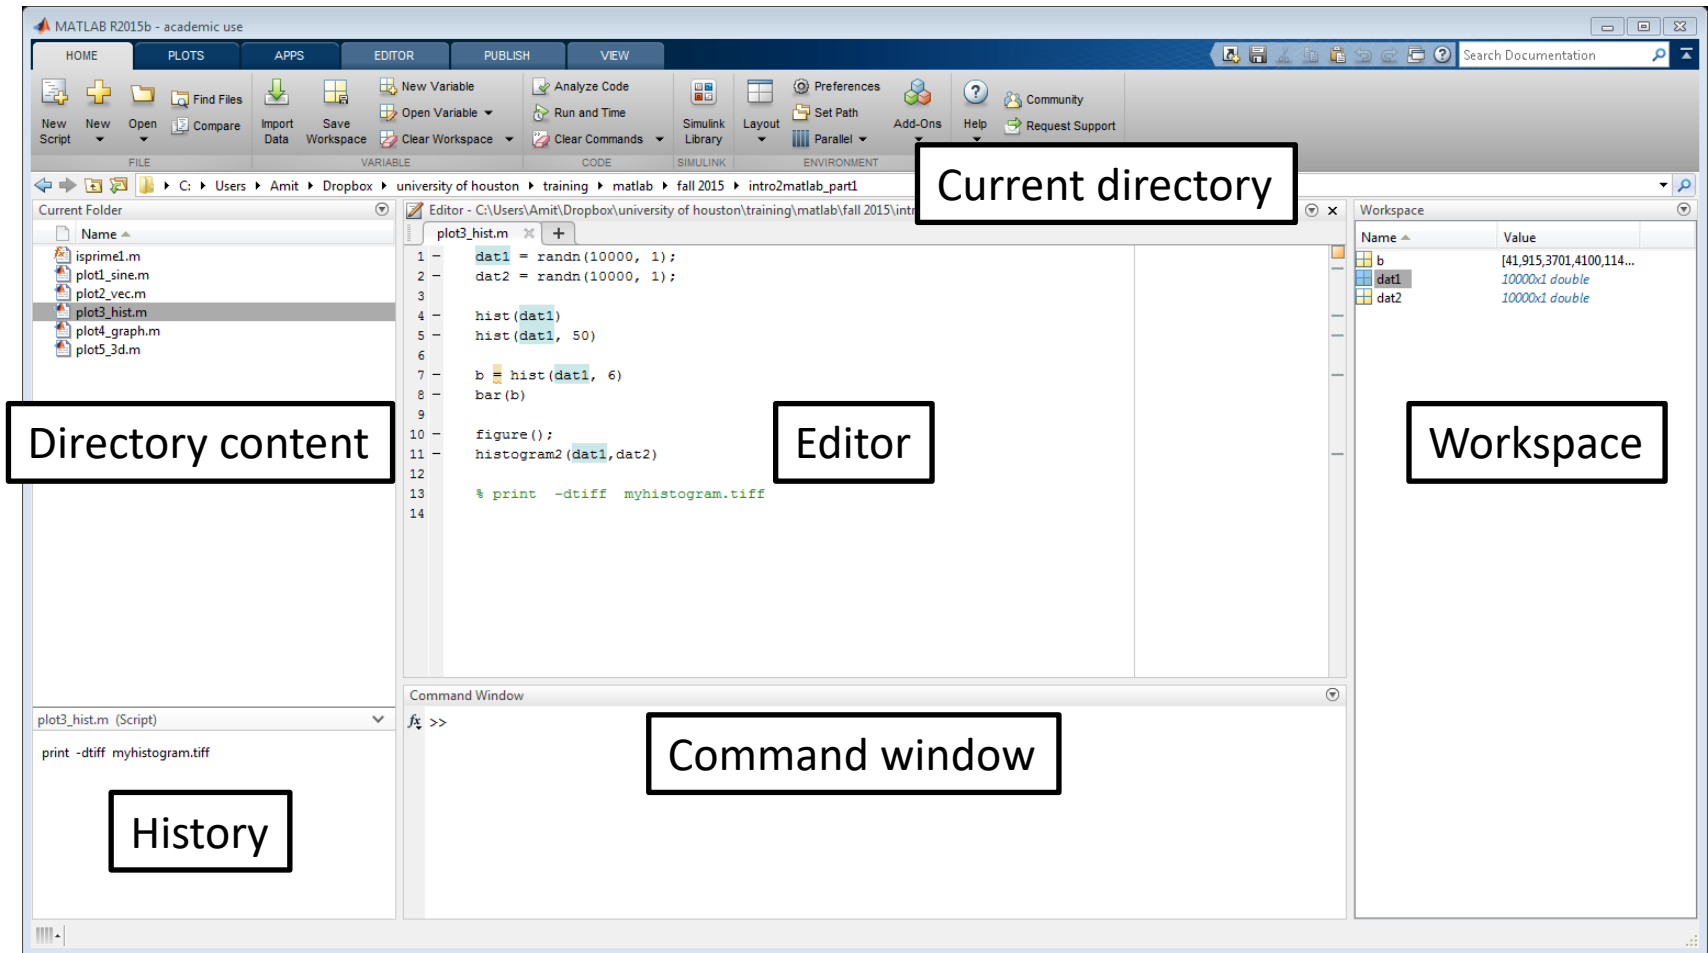
- Using M files

- Scripts and functions

- Plots




Review – MATLAB GUI



Review exercise

- Write a program to plot the Fibonacci series – first 20 numbers in the series
- Hint: use the fibonacci.m function file
- `>>plot(fibonacci(20));`



Review – parallel MATLAB

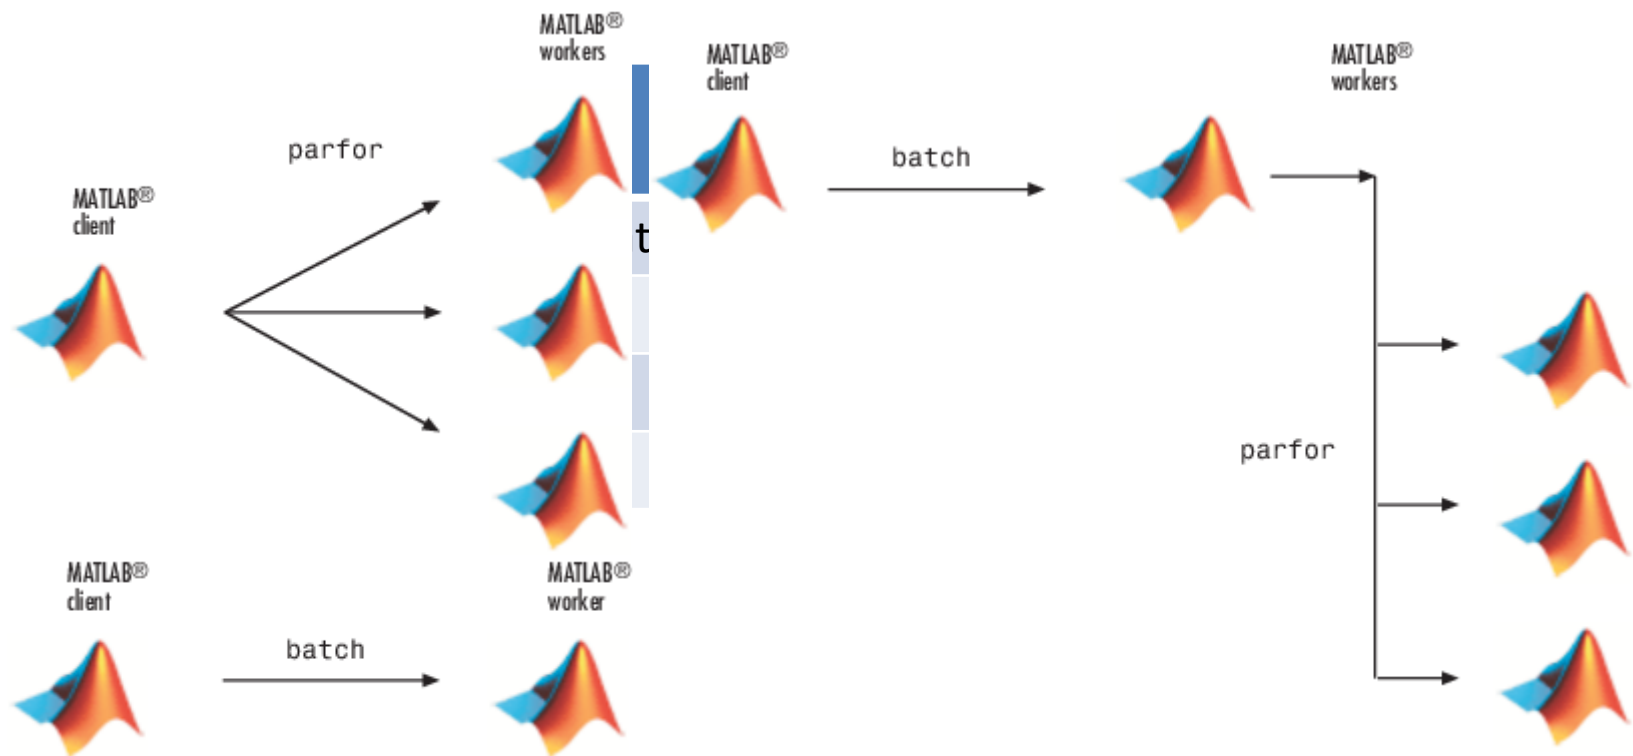
- Parallel MATLAB
 - extension of MATLAB that takes advantage of multicore desktop machines, GPUs and clusters.
- The Parallel Computing Toolbox or PCT runs on a desktop, and can take advantage of cores (R2014a has no limit, R2013b limit is 12, ...). Parallel programs can be run interactively or in batch.
- The MATLAB Distributed Computing Server (MDCS) controls parallel execution of MATLAB on a cluster with tens or hundreds of cores.

Review – parallel MATLAB

- Three ways to write a CPU parallel MATLAB program:
 - suitable for loops can be made into parfor loops;
 - the spmd statement can define cooperating synchronized processing;
 - the task feature creates multiple independent programs.
- GPU Computing
- The parfor approach is a limited but simple way to get started.
- spmd is powerful, but may require rethinking the program/data.
- The task approach is simple, but suitable only for computations that need almost no communication.

Review – parallel MATLAB

- There are several ways to execute a parallel MATLAB program:



Review exercise

- Parallelize the following function and measure its execution time for $n=10,000,000$

```
function total = prime_fun ( n )
%% PRIME returns the number of primes between 1 and N.
    total = 0 ;
    for i = 2 : n
        prime = 1 ;
        for j = 2 : i-1
            if ( mod ( i , j ) == 0 )
                prime = 0 ;
            end
        end
        total = total + prime ;
    end
    return
end
```

Use the prime_fun.m function file

Access Your Account

- UHVPN connection may be required if you are not on campus network
- Make sure that you are added to the classroom cluster access
- If you have confirmed enrollment then you should have access
- Ask the instructor if you have trouble

`ssh -XY -l username aerb202.cacds.e.uh.edu`

- Log into your accounts
- Username or login = CougarNet ID
- **Password = CougarNet password**

Interactive Job on GPU node

```
ssh -XY -l username aerb202.cacds.e.uh.edu
```

```
srun -N 1 -n 8 --x11 --pty /bin/bash
```

Interactive Job on GPU node

```
ssh -XY -l username aerb202.cacds.e.uh.edu
```

```
srun -N 1 -n 8 --x11 --pty /bin/bash
```

```
queue -u $USER
```

Accessing an Allocated GPU node

```
ssh -XY -l username aerb202.cacds.e.uh.edu
```

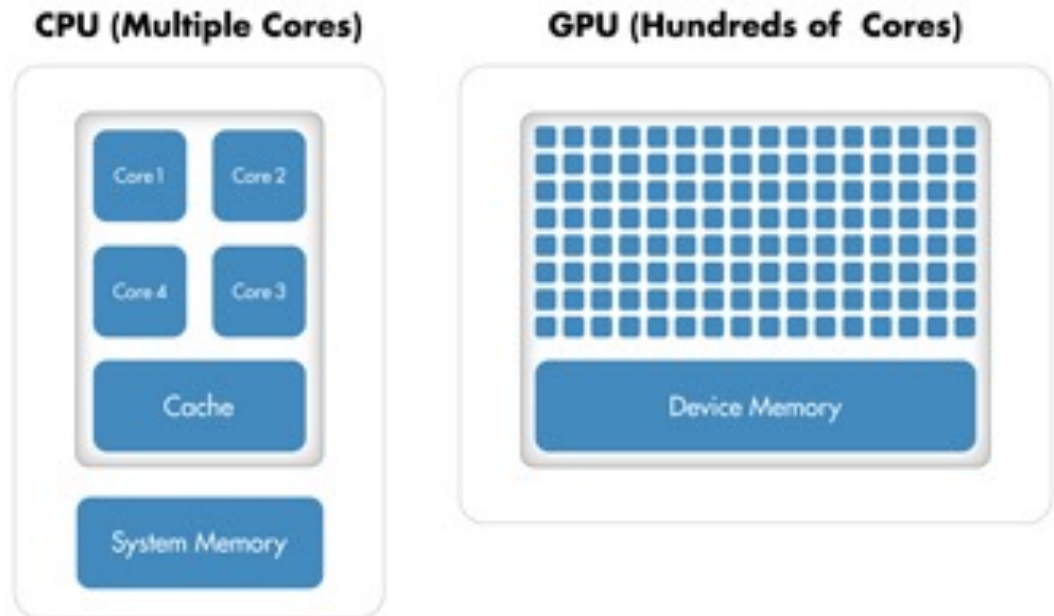
```
srun -N 1 -n 8 --x11 --pty /bin/bash
```

```
queue -u $USER
```

```
ssh -XY aerb-202?
```

GPU acceleration

- CPU
 - fast
 - general-purpose
- GPU
 - highly parallel
 - handles specific tasks with large amount of data
 - **memory transfers needed**



What/Why GPU computing?

- Serial portions of the code run on the CPU while parallel portions run on the GPU
- From a user's perspective, applications in general run significantly faster

Basic setup

- **CUDA-enabled NVIDIA GPUs** with compute capability 2.0 or higher. For releases 2014a and earlier, compute capability 1.3 is sufficient.
- Latest CUDA driver with 64 bit OS
- Your MATLAB must be version 2010b or later:
 - Go to the HELP menu, and choose About MATLAB.
- You must have the Parallel Computing Toolbox (PCT):
 - At UH, the concurrent (& student) license includes the PCT.
 - The standalone license does not include the PCT.
- To list all your toolboxes, type the MATLAB command **ver**.
- Limited availability on UH research cluster

Application requirement

- **Will Execution on a GPU Accelerate My Application?**
- A GPU can accelerate an application if it fits both of the following criteria:
 - **Computationally intensive**—The time spent on computation significantly exceeds the time spent on transferring data to and from GPU memory.
 - **Massively parallel**—The computations can be broken down into hundreds or thousands of independent units of work.
- Applications that do not satisfy these criteria might actually run slower on a GPU than on a CPU.

Basic usage

- Send data to GPU
 - either allocate there or transfer from workspace
- Run Matlab functions
 - GPU acceleration is used automatically
 - Built-in and custom functions
- Retrieve the output data

GPU Capabilities and Performance

- Identify and Select a GPU Device
- Transfer or Create arrays on GPU
- Run built-in functions on GPU
- Run Element-wise MATLAB Code on GPU
- Run CUDA or PTX Code on GPU
- Run MEX-Functions Containing CUDA Code

Identify and Select a GPU Device

- GPU functions
 - compute capability 2.0 and above

<code>gpuDevice</code>	Query or select GPU device
<code>gpuDeviceCount</code>	Number of GPU devices present
<code>gputimeit</code>	Time required to run function on GPU
<code>reset</code>	Reset GPU device and clear its memory
<code>wait (GPUDevice)</code>	Wait for GPU calculation to complete

>> `gpuDeviceCount`

>> `gpuDevice`

Transfer Arrays Between Workspace and GPU

- Transfer Arrays Between Workspace and GPU
 - `X = rand(100);`
 - `G = gpuArray(X);`
- Create GPU Arrays Directly
 - `G = rand(100,'gpuArray');`
- Transfer data back to MATLAB workspace
 - `C = gather(G);`

Try these examples:

Create on GPUs

```
>> identity = eye(1024,'int32','gpuArray');
```

```
>> Z = zeros(8192,1,'gpuArray');
```

Send to GPUs

```
>> N = 6;
```

```
>> M = magic(N);
```

```
>> G = gpuArray(M);
```

```
>> G = gpuArray(single(X));
```

```
>> G = gpuArray(ones(100,'uint32'));
```

Retrieve from GPUs

```
>> G = gpuArray(ones(100,'uint32')); %Array stored on GPU
```

```
>> D = gather(G); % Get G from GPU to D in MATLAB workspace
```

```
>> OK = isequal(D,ones(100,'uint32')) % check if G on GPU is same as D  
on CPU
```

Examine gpuArray Characteristics

Function	Description
<code>classUnderlying</code>	Class of the underlying data in the array
<code>existsOnGPU</code>	Indication if array exists on the GPU and is accessible
<code>isreal</code>	Indication if array data is real
<code>length</code>	Length of vector or largest array dimension
<code>ndims</code>	Number of dimensions in the array
<code>size</code>	Size of array dimensions

GPUArray class

`gpuArray`

- main data class for GPU computations
- stored in the GPU memory
- create directly using static methods

<code>zeros</code>	<code>nan</code>	<code>eye</code>	<code>rand</code>	<code>linspace</code>
<code>ones</code>	<code>true</code>	<code>colon</code>	<code>randi</code>	<code>logspace</code>
<code>inf</code>	<code>false</code>		<code>randn</code>	

- copy from existing data

`gpuArray(img)`

GPUArray class

- Supported data types:
(u)int8, (u)int16, (u)int32, (u)int64, single, double, logical
 - determine the type using
`classUnderlying(gpuVar)`
- Retrieve the data using
`workspaceVar = gather(gpuVar)`

Run Built-In Functions on a GPU

- If at least one of the input arguments is a `gpuArray`, the function executes on the GPU and returns a `gpuArray`
- Common functions on GPUs - discrete Fourier transform (`fft`), matrix multiplication (`mtimes`), left matrix division (`mldivide`)
- Full list - `methods('gpuArray')` or <http://www.mathworks.com/help/distcomp/run-built-in-functions-on-a-gpu.html>
- `help gpuArray/functionname`
- Explicit definition for Complex numbers if output is expected to be complex
 - `G = gpuArray(complex(p))`

Functions with gpuArray

Compare the execution time of the fft function on GPU vs CPU

gfun.m

```
Ga = rand(1000,'single','gpuArray');  
Gfft = fft(Ga);  
Gb = (real(Gfft) + Ga) * 6;  
G = gather(Gb);  
whos
```

```
Ga = rand(1000,'single');  
Gfft = fft(Ga);  
Gb = (real(Gfft) + Ga) * 6;  
whos
```

cfun.m

Sparse matrices on GPUs

sparsefun.m

```
x = [0 1 0 0 0; 0 0 0 0 1]  
s = sparse(x)  
g = gpuArray(s); % g is a sparse gpuArray  
gt = transpose(g); % gt is a sparse gpuArray  
f = full(gt)  
whos
```

Simple example

- Solve system of linear equations ($Ax = b$)

```
A = gpuArray(A) ;
```

```
b = gpuArray(b) ;
```

```
x = A\b ;
```

```
x = gather(x) ;
```

Simple example

- Compute convolution using FFT (convolution.m)

```
img = gpuArray(img);  
msk = padarray(msk,size(img)-size(msk),0,'post');  
msk = gpuArray(msk);  
I = fft2(img);  
M = fft2(msk,size(img,1),size(img,2));  
res = real(ifft2(I.*M));  
res = gather(res);
```

Run Element-wise MATLAB Code on GPU

- You can run your own MATLAB function of element-wise operations on a GPU

- Define arrays on GPU
- Use `arrayfun` or `bsxfun` to execute the custom functions

- List of supported MATLAB

- <http://www.mathworks.com/help/gpu/element-wise-matlab-code-gpu.html#bsnx7h8-1>

```
function Y = myfun(X)
    R = rand();
    Y = R.*X;
end
G = 2*ones(4,4,'gpuArray')
H = arrayfun(@myfun, G)
```

Random numbers on GPUs


- rand, randi, and randn partially available on GPUs
- Default random stream on CPU and GPU are not same
 - parallel.gpu.rng
 - parallel.gpu.RandStream
- To set the streams equal
 - Set the same stream and seed
 - Use the random_equal.m

MEX-Functions with CUDA Code

- Write a MEX-File Containing CUDA Code
- Compile a GPU MEX-File
- Run the Resulting MEX-Functions

Benchmarking

- Run the MATLAB benchmark for GPU
 - Use script `gpu_benchmark.m`
- Run the MATLAB benchmark for Host
 - Use script `gpu_benchmark.m`



Exercise - The Mandelbrot Set

- Using CPU only
 - mandelbrot_cpu.m
- Using the existing algorithm but with GPU data as input
 - mandelbrot_gpuArray.m
- Using arrayfun to perform the algorithm on each element independently
 - mandelbrot_arrayfun.m
- *Using the MATLAB/CUDA interface to run some existing CUDA/C++ code*

GPU Code porting

- Profile the existing code
 - Identify the hotspots using MTLAB profile tool
 - Custom functions
 - Parfor loops
 - Custom MEX functions
 - Built-in functions without corresponding gpuArray version
 - High frequency of function calls for many functions
 - Short execution time
 - Use the MATLAB debugger to get memory requirements
 - Large input/output data
 - Simple data types (arrays and matrices)

Code porting strategies

- Use the built-in gupArray functions
 - Check the results for correctness
 - Speedup might vary
 - Faster than CPU – good
 - Slower than CPU – not bad!
 - At least it works and other optimizations might improve performance
- Matrix manipulations tend to be faster on GPUs
 - Move operations like reshape, subsref, etc...

Code porting strategies

- Eliminate loops/calls in the code
 - Vectorize the code

```
A = rand(4);  
output = zeros(4);  
for n=1:4  
    output(n,:) = input(n,:) / n;  
end
```



```
A=rand(4);  
output=zeros(4);  
output = A ./ repmat([1:4].',[1,4]);
```

- Minimize the calls to gpuArray functions to avoid data transfers

Code porting strategies

- Wisely use the limited GPU memory
 - Keep reuse data on GPUs

??? Error using ==> gpuArray at 37

Out of memory on device. You requested: 651.73Mb, device has 1.27Gb free.

- use “clear” to flush unnecessary variables from GPU memory *or*
- divide your problem into smaller chunks
- Use host memory as buffer for data transfer

Code porting strategies

- Recast for arrayfun
 - element-wise operations
 - operations across a large number of scalar values
 - handle bulk processing of small arrays

```
R1 = rand(2,5,4,'gpuArray');  
R2 = rand(2,1,4,3,'gpuArray');  
R3 = rand(1,5,4,3,'gpuArray');  
R = arrayfun(@(x,y,z)(x+y.*z),R1,R2,R3);  
size(R)
```

- Combine vectorized statements

```
>> result = (a1 + a2 + a3) ./ 3;
```



```
function out=littleaaafun(a1,a2,a3)  
out = (a1+a2+a3) / 3;
```


GPU arrayfun example

- Define a MATLAB function

```
function [o1,o2] = aGpuFunction(a,b,c)
o1 = a + b;
o2 = o1 .* c + 2;
```

- Evaluate on GPU

```
s1 = gpuArray(rand(400));
s2 = gpuArray(rand(400));
s3 = gpuArray(rand(400));
[o1,o2] = arrayfun(@aGpuFunction,s1,s2,s3);
whos
```

- Retrieve the data to MATLAB workspace on CPU

```
d = gather(o2);
```

Code porting strategies

- Create a CUDA kernel
 - big speed-up and a big headache
 - Simple functions that are task independent
 - translate an existing MEX function that is very simple to run on the GPU
 - complex code that include a lot of branching, task dependencies, and/or serialized output

Code porting strategies for multiple GPUs

- Use GPUs from a parpool on CPUs
 - GPU function call from a parfor could create conflicts in worker processes so GPU computations in parfor are serialized across workers.
 - Use spmd to pin a GPU per worker process

```
parpool('local',4)
spmd
    gpuDevice( labindex );
    % customer GPU code goes here
end
delete(gcp)
```

- Alternatively, toggle between CPU and GPU
 - See code in testgpu.m

Code porting strategies

- Test code changes
 - Use 'whos' to see if GPUs are actually being used
 - Due to changes in handling of arithmetic on CPU and GPU a small amount of variance in results (difference of $O(1e-7)$ percent may be reasonable)
- Find and fix new bottlenecks
 - Incrementally add GPU parallelism
 - Scale up the input data size to find other issues
 - Iteratively fix the bottlenecks
 - Rethink the algorithm if necessary

References

- <http://www.mathworks.com/discovery/matlab-gpu.html>
- <http://www.mathworks.com/company/newsletters/articles/gpu-programming-in-matlab.html>
- http://its2.unc.edu/divisions/rc/training/scientific/short_courses/ParallelGPUMATLAB.pptx
- <http://zoi.utia.cas.cz/files/GPU%20acceleration%20in%20Matlab.pptx>
- <http://www.cac.cornell.edu/matlab/TechDocs/Examples/BestPracticesGPU.aspx>
- <http://www.mathworks.com/moler/exm/book.pdf>